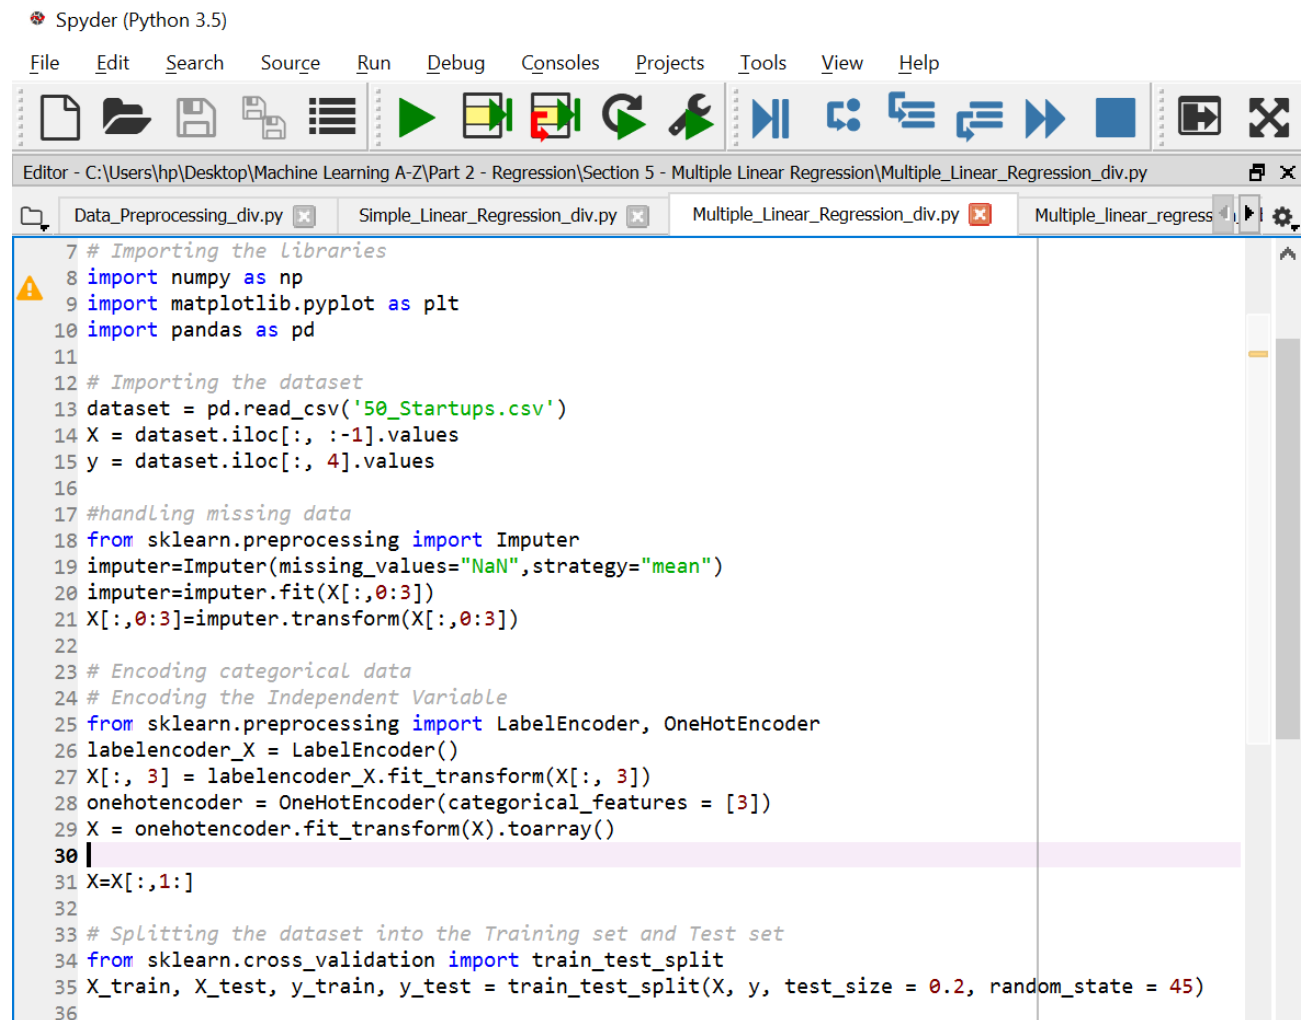


# Implementation of Multiple Linear Regression using Backward Elimination

CODE::



The image shows the Spyder Python IDE interface. The top menu bar includes File, Edit, Search, Source, Run, Debug, Consoles, Projects, Tools, View, and Help. Below the menu is a toolbar with various icons for file operations and execution. The editor window displays a Python script for Multiple Linear Regression using Backward Elimination. The script imports necessary libraries, loads a dataset, handles missing data, encodes categorical data, and splits the dataset into training and testing sets.

```
7 # Importing the libraries
8 import numpy as np
9 import matplotlib.pyplot as plt
10 import pandas as pd
11
12 # Importing the dataset
13 dataset = pd.read_csv('50_Startups.csv')
14 X = dataset.iloc[:, :-1].values
15 y = dataset.iloc[:, 4].values
16
17 #handling missing data
18 from sklearn.preprocessing import Imputer
19 imputer=Imputer(missing_values="NaN",strategy="mean")
20 imputer=imputer.fit(X[:,0:3])
21 X[:,0:3]=imputer.transform(X[:,0:3])
22
23 # Encoding categorical data
24 # Encoding the Independent Variable
25 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
26 labelencoder_X = LabelEncoder()
27 X[:, 3] = labelencoder_X.fit_transform(X[:, 3])
28 onehotencoder = OneHotEncoder(categorical_features = [3])
29 X = onehotencoder.fit_transform(X).toarray()
30
31 X=X[:,1:]
32
33 # Splitting the dataset into the Training set and Test set
34 from sklearn.cross_validation import train_test_split
35 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 45)
36
```

Spyder (Python 3.5)

File Edit Search Source Run Debug Consoles Projects Tools View Help



Editor - C:\Users\hp\Desktop\Machine Learning A-Z\Part 2 - Regression\Section 5 - Multiple Linear Regression\Multiple\_Linear\_Regression\_div.py

Data\_Preprocessing\_div.py Simple\_Linear\_Regression\_div.py Multiple\_Linear\_Regression\_div.py Multiple\_linear\_regres

```
37 # Feature Scaling
38 """from sklearn.preprocessing import StandardScaler
39 sc_X = StandardScaler()
40 X_train = sc_X.fit_transform(X_train)
41 X_test = sc_X.transform(X_test)
42 sc_y = StandardScaler()
43 y_train = sc_y.fit_transform(y_train)"""
44
45 #fitting multiple linear regression to training set
46 from sklearn.linear_model import LinearRegression
47 regressor=LinearRegression()
48 regressor.fit(X_train,y_train)
49
50 #predicting the test set results
51 y_pred=regressor.predict(X_test)
52
53 #building optimal solution using backward elimination
54 import statsmodels.formula.api as sm
55 X=np.append(arr=np.ones((50,1)).astype(int),values=X,axis=1)
56 X_opt=X[:,[0,1,2,3,4,5]]
57 regressor_OLS=sm.OLS(endog=y,exog=X_opt).fit()
58 regressor_OLS.summary()
59 X_opt=X[:,[0,1,3,4,5]]
60 regressor_OLS=sm.OLS(endog=y,exog=X_opt).fit()
61 regressor_OLS.summary()
62 X_opt=X[:,[0,3,4,5]]
63 regressor_OLS=sm.OLS(endog=y,exog=X_opt).fit()
64 regressor_OLS.summary()
65 X_opt=X[:,[0,3,5]]
66 regressor_OLS=sm.OLS(endog=y,exog=X_opt).fit()
67 regressor_OLS.summary()
68 X_opt=X[:,[0,3]]
69 regressor_OLS=sm.OLS(endog=y,exog=X_opt).fit()
70 regressor_OLS.summary()
```

Permissions: RW

End-of-lines:

dataset - DataFrame

Index	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349	136898	471784	New York	192262
1	162598	151378	443899	California	191792
2	153442	101146	407935	Florida	191050
3	144372	118672	383200	New York	182902
4	142107	91391.8	366168	Florida	166188
5	131877	99814.7	362861	New York	156991
6	134615	147199	127717	California	156123
7	130298	145530	323877	Florida	155753
8	120543	148719	311613	New York	152212
9	123335	108679	304982	California	149760
10	101913	110594	229161	Florida	146122
11	100672	91790.6	249745	California	144259
12	93863.8	127320	249839	Florida	141586
13	91992.4	135495	252665	California	134307

Format    Resize    ☒ Background color    ☒ Column min/max    Save and Close

Fig1:: dataset=pd.read\_csv('50\_Startups.csv')

```
[[165349.2 136897.8 471784.1 'New York']
 [162597.7 151377.59 443898.53 'California']
 [153441.51 101145.55 407934.54 'Florida']
 [144372.41 118671.85 383199.62 'New York']
 [142107.34 91391.77 366168.42 'Florida']
 [131876.9 99814.71 362861.36 'New York']
 [134615.46 147198.87 127716.82 'California']
 [130298.13 145530.06 323876.68 'Florida']
 [120542.52 148718.95 311613.29 'New York']
 [123334.88 108679.17 304981.62 'California']
 [101913.08 110594.11 229160.95 'Florida']
 [100671.96 91790.61 249744.55 'California']
 [93863.75 127320.38 249839.44 'Florida']
 [91992.39 135495.07 252664.93 'California']
 [119943.24 156547.42 256512.92 'Florida']
 [114523.61 122616.84 261776.23 'New York']
 [78013.11 121597.55 264346.06 'California']
 [94657.16 145077.58 282574.31 'New York']
 [91749.16 114175.79 294919.57 'Florida']
 [86419.7 153514.11 0.0 'New York']
 [76253.86 113867.3 298664.47 'California']
 [78389.47 153773.43 299737.29 'New York']
 [73994.56 122782.75 303319.26 'Florida']
 [67532.53 105751.03 304768.73 'Florida']
 [77044.01 99281.34 140574.81 'New York']
 [64664.71 139553.16 137962.62 'California']
 [75328.87 144135.98 134050.07 'Florida']
 ... ..]
```

Fig2:: x=dataset.iloc[:,0:4].values

y - NumPy array

	0
0	192262
1	191792
2	191050
3	182902
4	166188
5	156991
6	156123
7	155753
8	152212
9	149760
10	146122
11	144259
12	141586

Fig3:: y=dataset.iloc[:,4].values

```
[165349.2 136897.8 471784.1 2]
[162597.7 151377.59 443898.53 0]
[153441.51 101145.55 407934.54 1]
[144372.41 118671.85 383199.62 2]
[142107.34 91391.77 366168.42 1]
[131876.9 99814.71 362861.36 2]
[134615.46 147198.87 127716.82 0]
[130298.13 145530.06 323876.68 1]
[120542.52 148718.95 311613.29 2]
[123334.88 108679.17 304981.62 0]
[101913.08 110594.11 229160.95 1]
[100671.96 91790.61 249744.55 0]
[93863.75 127320.38 249839.44 1]
[91992.39 135495.07 252664.93 0]
[119943.24 156547.42 256512.92 1]
[114523.61 122616.84 261776.23 2]
[78013.11 121597.55 264346.06 0]
[94657.16 145077.58 282574.31 2]
[91749.16 114175.79 294919.57 1]
[86419.7 153514.11 0.0 2]
[76253.86 113867.3 298664.47 0]
[78389.47 153773.43 299737.29 2]
[73994.56 122782.75 303319.26 1]
[67532.53 105751.03 304768.73 1]
[77044.01 99281.34 140574.81 2]
[64664.71 139553.16 137962.62 0]
```

Fig4:: x[:,3]=labelencoder.fit\_transform(x[:,3])

x - NumPy array

	0	1	2	3	4	5
0	0	0	1	165349	136898	471784
1	1	0	0	162598	151378	443899
2	0	1	0	153442	101146	407935
3	0	0	1	144372	118672	383200
4	0	1	0	142107	91391.8	366168
5	0	0	1	131877	99814.7	362861
6	1	0	0	134615	147199	127717
7	0	1	0	130298	145530	323877
8	0	0	1	120543	148719	311613
9	1	0	0	123335	108679	304982
10	0	1	0	101913	110594	229161
11	1	0	0	100672	91790.6	249745
12	0	1	0	93863.8	127320	249839

Format Resize ☒ Background color

Fig5:: x=onehotencoder.fit\_transform(x).toarray()

x - NumPy array

	0	1	2	3	4
0	0	1	165349	136898	471784
1	0	0	162598	151378	443899
2	1	0	153442	101146	407935
3	0	1	144372	118672	383200
4	1	0	142107	91391.8	366168
5	0	1	131877	99814.7	362861
6	0	0	134615	147199	127717
7	1	0	130298	145530	323877
8	0	1	120543	148719	311613
9	0	0	123335	108679	304982
10	1	0	101913	110594	229161
11	0	0	100672	91790.6	249745
12	1	0	93863.8	127320	249839

Fig6:: x=x[:,1:]

Dataset after the removal of dummy variable